



Reward Distribution Analysis


Overall Reward Distribution




Reward Distribution by Pod



TTFT vs TPOT (colored by Reward)







Pod Performance Comparison







Feature Probability Distributions


Cpu Kv Cache
 $\mu=0.00$, $\sigma=0.00$




Kv Hit Ratio
 $\mu=3.67$, $\sigma=4.82$




Input Tokens
 $\mu=3897.40$, $\sigma=2179.75$




Output Tokens
 $\mu=100.00$, $\sigma=0.00$





Total Tokens
 $\mu=3997.40$, $\sigma=2179.75$



Ttft Reward Weight Used
 $\mu=1.00$, $\sigma=0.00$




State-Performance Correlation Analysis




Reward Variance & Model Confidence Prediction Analysis


Reward Variance vs Known Benchmarks



Reward Gap vs Known Benchmarks



Confidence Calibration Prediction Matrix



Dataset Quality Summary

