



Reward Distribution Analysis


Overall Reward Distribution




Reward Distribution by Pod



TTFT vs TPOT (colored by Reward)







Pod Performance Comparison







Feature Probability Distributions


Cpu Kv Cache
 $\mu=0.00$, $\sigma=0.00$




Kv Hit Ratio
 $\mu=17.26$, $\sigma=32.81$




Input Tokens
 $\mu=5867.42$, $\sigma=2972.31$




Output Tokens
 $\mu=100.00$, $\sigma=0.00$






Total Tokens
 $\mu=5967.42$, $\sigma=2972.31$



Ttft Reward Weight Used
 $\mu=2.00$, $\sigma=0.00$




State-Performance Correlation Analysis




Reward Variance & Model Confidence Prediction Analysis


Reward Variance vs Known Benchmarks



Reward Gap vs Known Benchmarks



Confidence Calibration Prediction Matrix



Dataset Quality Summary

