



Reward Distribution Analysis


Overall Reward Distribution




Reward Distribution by Pod



TTFT vs TPOT (colored by Reward)







Pod Performance Comparison







Feature Probability Distributions


Cpu Kv Cache
 $\mu=0.00$, $\sigma=0.00$




Kv Hit Ratio
 $\mu=16.68$, $\sigma=32.40$




Input Tokens
 $\mu=6037.15$, $\sigma=3025.51$





Output Tokens
 $\mu=100.00$, $\sigma=0.02$



Total Tokens
 $\mu=6137.15$, $\sigma=3025.51$




State-Performance Correlation Analysis




Reward Variance & Model Confidence Prediction Analysis


Reward Variance vs Known Benchmarks



Reward Gap vs Known Benchmarks



Confidence Calibration Prediction Matrix



Dataset Quality Summary

